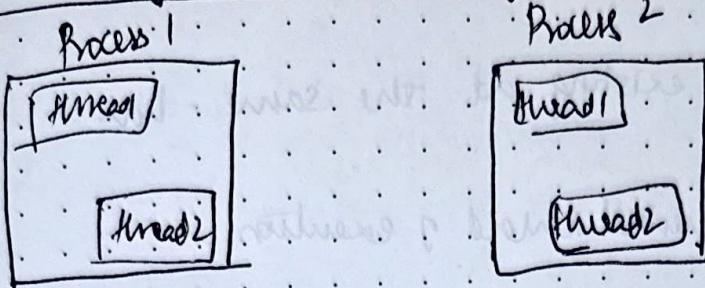


## Multithreading and Concurrency



Process : Instance of a program that is getting executed

It has its own resource like memory, thread etc. OS allocates resources to process when it's created

(javac Test.java)

compilation → generates bytecode that can be executed by JVM

↳ execution → at this point, JVM starts the new process, here the class which has "public static void main (String args)" method

Each process has its own main-memory

Thread : lightweight process

or

smallest sequence of instructions that are executed by CPU independently

→ 1. process can have multiple threads

→ when a process is created, it starts with 1 thread and that initial 'main thread' and from that we can create multiple threads to perform task concurrently

## Concurrency

→ at (con.) more events happening or existing at the same time

Typically, we have one main sequential thread of execution

One CPU core executes serially

But what if we could have a second path in our code executing

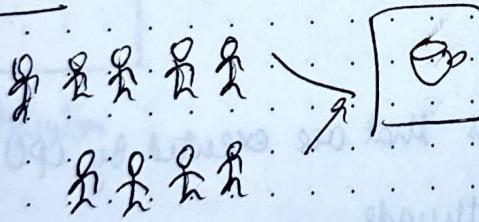
"Performance is the currency of computing"

## Parallelism vs Concurrency



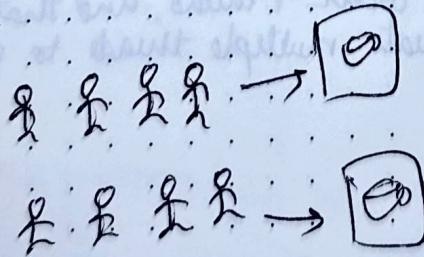
- multiple things can happen at once, the order matters, and sometimes tasks have to wait on shared resources
- Parallelism definition: Everything happens at once, instantaneously

### concurrent



coffee machine

### Parallel



## Necessity of concurrency

- Concurrency (like parallelism) is essential because it is necessary for a system to function.

## Thread based concurrency using C++

main thread

main

foo()

bar()

baz()

Thread 1

func()

func2()

2 paths

- 1 process has multiple threads

- Each thread shares the same code, data & kernel context

- Thread has its own TID

- Thread has its own logical control flow

- Thread has its own stack for local variables

## When to use threads

- When we have heavy computation

- Using threads on your GPU for graphics

- Use threads to separate work

- Gives performance

- Also simplifies the logic of your problem

## → Standard C++ thread library

(a) pthread by Intel

```
#include <iostream>
#include <thread>
```

```
int main() {
    std::cout << "Hello" << std::endl; // Single threaded
    return 0;
}
```

By adding one more fn.

```
void test(int x) {
    std::cout << "Hello from thread!" << std::endl;
    std::cout << "Argument:" << x << std::endl;
```

```
int main().
```

```
std::thread myThread(&test, 100);
myThread.join();
```

→ We need to make the main thread wait, so that it executes in order

join() → blocks the current thread until the thread identified by  
this finishes its execution.

→ Creating threads using lambda function.

Launching a multiple threads from a vector

```
int main() {
    auto lambda = [](int x) {
        std::cout << "Hello from thread" << std::this_thread::get_id() <<
        std::endl;
    };
    std::cout << "Argument passed in:" << x << std::endl;
    std::vector<std::thread> threads;
    for(int i=0; i<10; i++) {
        threads.push_back(std::thread(lambda, i));
        threads[i].join();
    }
    std::cout << "Hello from my thread" << std::endl;
    return 0;
}
```

But getting same thread id, why?

so it is nothing but a sequential code.

so break & put this into another for loop.

Now it is correct.

but the output is jumbled up.

Interleaving of threads is happening, so how to handle this synchronization.

## Introducing the `gthread`

`gthread` for support with auto-joining and cancellation.

C++20: PAAJ → Resource acquisition is initialisation

using `gthread`, we do not have to do the `join()`.

Joins at the end of the scope

How to have threads working on a shared problem?

How to share data & improve our performance

static int shared\_value = 0;

void shared\_v.increment() {

    shared\_value = shared\_value + 1;

}

int main() {

    std::vector<std::thread> threads;

    for (int i = 0; i < 10; i++) {

        threads.push\_back(std::thread(shared\_value, increment));

}

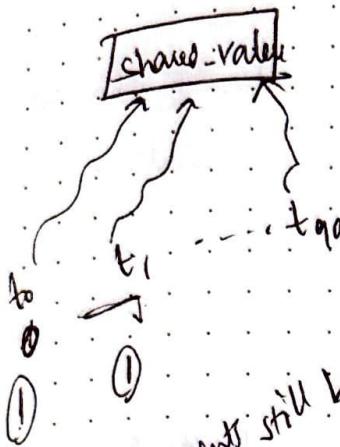
    for (int i = 0; i < 10; i++) {

        threads[i].join();

}

    std::cout << "shared value" << shared\_value << std::endl;

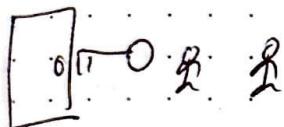
}



This is called data race condition.

So, how to fix this?

mutex → mutual exclusion



threads access shared state. then read/write  
only after its done. - thread 1 does it.

#include <mutex>

std::mutex glock;

→ now inside this

void shared\_value\_increment() {

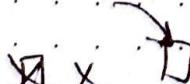
    glock.lock();

    shared\_value += 1;

    glock.unlock();

}

$t_0, t_1, t_2, t_3, \dots, t_{q4}$



shared-value

lock() → locks the mutex, blocks if the mutex is not available

unlock() → unlocks the mutex

Now, it works.

say  $t_2$  acquires the lock first,  
then all others are blocked.

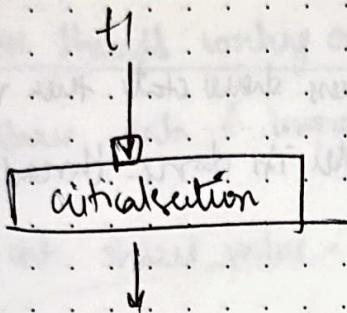
\* Want this slow down the performance?

No, critical sections should be protected

→ Thread ~~scheduler~~ creation is non-deterministic, and execution of threads is not sequential. It is determined by the OS thread scheduler.

### Preventing deadlock:

say we forgot to unlock the code, and just one program holds the lock and so the program gets freezed.



but doesn't return the lock.

Deadlock → blocked (n)

say      lock();  
try {

    critical section

    }

    but some exception is thrown;

    except {

    }

}

    }

    unlock();

}

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

    }

lock\_guard() → is a mutex wrapper.  
The class lock\_guard is a mutex wrapper that provides a RAII-style mechanism for own duration of scoped block.  
When a lock\_guard is created, it attempts to take ownership of the mutex it is given.

std::lock\_guard<std::mutex> lockguard(glock);

void shared\_value\_inrement() {

std::lock\_guard<std::mutex> lockguard(gLock);

shared\_value = shared\_value + 1;

}

scoped\_lock(C++17)

lock\_guard(C++11) → Instead of manually locking & unlocking,

Even in presence of try & catch, this works.

Another way to update shared\_value without lock

std::atomic

#include<atomic>

can get rid of our locks.

static std::atomic<int> shared\_value = 0;

void shared\_value\_inrement() {

shared\_value ++;

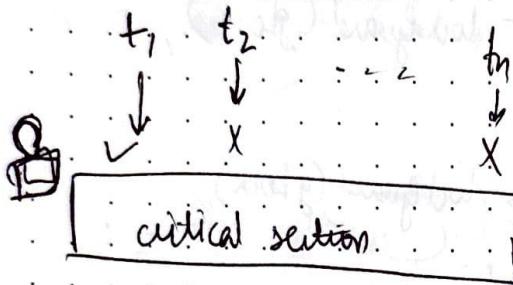
}

Any data type are able to do this

operator ++      operator +=      operator &=  
operator --      operator -=      operator |=

use only  
here  $\Rightarrow$

### Synchronization of threads with locks



→ What does other threads do when  $t_1$  is accessing the critical section?

They are waiting, happens when there is spin lock, as this results in wastage of CPU cycles as they are waiting, because other threads are constantly asking, do I have the lock on?

There must be more efficient way

`std::condition_variable`

↳ synchronous primitive that can be used to block a thread  
(n) multiple threads at the same time until another thread modifies a shared variable, and notifies the condition variable

→ event, notify, notify\_all on the `std::condition_variable`

1. boolean lock - unique lock
2. condition-variable
3. 2 threads, worker / reporter

```

std::mutex glock;
std::condition_variable gConditionVariable;
int main() {
    result = 0;
    bool notified = false;
    // Reporting thread, // Must wait on work done by the working thread
    std::thread reporter([&] {
        gConditionVariable.wait(glock);
    });
    // working thread
    std::thread worker([&] {
        std::unique_lock<std::mutex> lock(glock);
        result = 12 + 1 + 7;
        // Our work is done
        notified = true;
        std::this_thread::sleep_for(std::chrono::seconds(1));
    });
    reporter.join();
    if (notified) {
        std::cout << "Result is " << result << std::endl;
    }
}

```

unique\_lock → more power than the lock-guard.

\* triangle

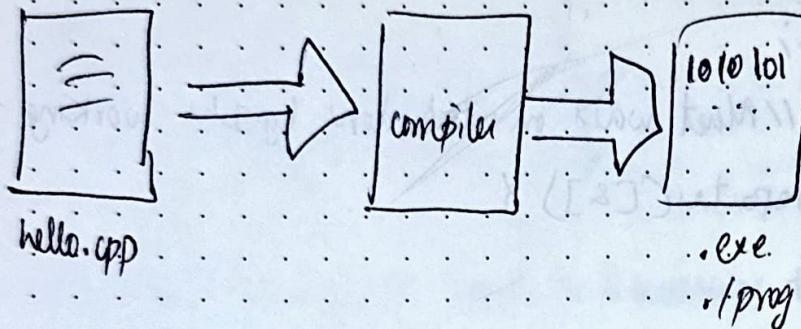
std::async

↳ without synchronization

Launch a thread & won't be blocked on that thread

Just like node.js' async keyword

Compiler



segmentation fault  $\rightarrow$  stack overflow

~~#~~ address of operator

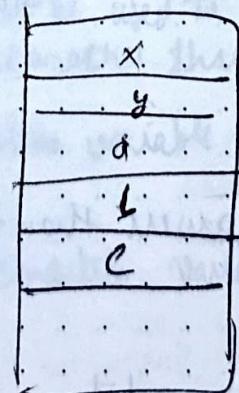
where actually in memory these variables live

#include <iostream>

int main()

```
int x=42;
float y=72;
char a='a';
signed char b='b';
unsigned char c='c';
return 0;
```

stack



y  
ptr

std::cout << "x & y" << endl;

Address is returned in x&y in hexadecinal.

but

$\text{sizeof}(x) \rightarrow 4 \text{ bytes}$  (to store an integer)

y → next location after x, after 4 bytes

a: abc<sup>a</sup>

b: bc<sup>a</sup>

c: c<sup>a</sup>

strange right?

(void\*) &a

now gives proper

(void\*) &b

address

(void\*) &c

They are offset only by 1 byte, as  
char takes only 1 byte.

void foo()

o/p : 1 , strange

y

&foo

→ addresses the function

so we do (void\*) &foo

You can think of & as a function itself.

Pass By Value

"arguments"

int main()

int x = 42

std::cout << "Value: " << x << "X address" << &x << std::endl;

setValue(x)

std::cout <<

return 0;

void setValue(int arg){

arg = arg; } → std::cout << arg << endl;

y → std::cout << arg << endl;

11

y

## Output

X value 42, address 0x7ff...b7h

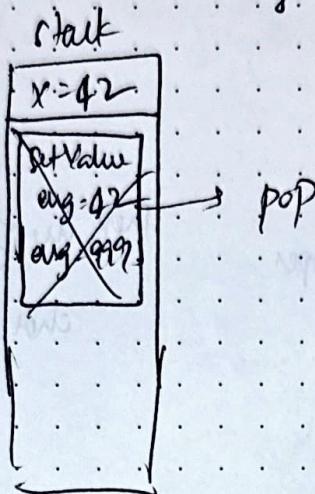
setValue: arg is initially 42

setValue: arg at end is 999.

X value 42, address same

"Make copies of our arguments when passed 'in function'"

But arg address is diff than X address



## References

2<sup>nd</sup> use of "&" symbol

ALIAS

int& ref = x;

// "int&" is a full type for a reference type

cout << ref

cout << &ref

both x & ref give same value &  
same reference

→ Another variable name that can refer to x.

→ Same named location memory as x.

Common use → pass by reference

```
#include <typeinfo>
typeid(x).name()
typeid(y).name()
```

$x\text{ type} \rightarrow i$   
 $y\text{ type} \rightarrow j$

## Pass By Reference

- 1) Mutate Data
- 2) Efficient

```
void PassByReference (int& arg)
arg = 999;
y
```

$\times$  address remains same  
 $\times$  value changes to 999.

```
n = 12
cout << n
cout << &n
PassByReference(x),
cout << n
cout << &n
```

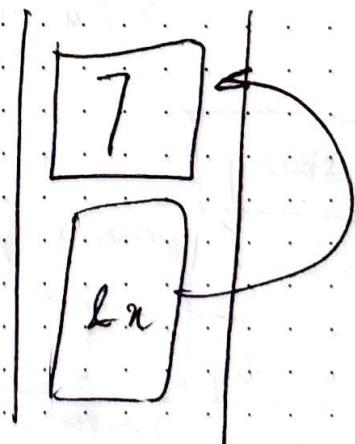
We avoided making a copy &  
mutating the same variable

```
#include <iostream>
```

```
#include <algorithm>
```

Pointers  $\rightarrow$  A datatype that stores an address:

```
int x = 7
int* px = &x;
cout << px;
```



What are we going to do with this value?

## Dereferencing the value

Retain the value / type of what we point to

"px Dereferenced" `<< *px << std::endl;`

int main () {

int x = 7 ;

int \*px = &x;

\*px = 42 ;

cout << x

o/p is 42

HEAP

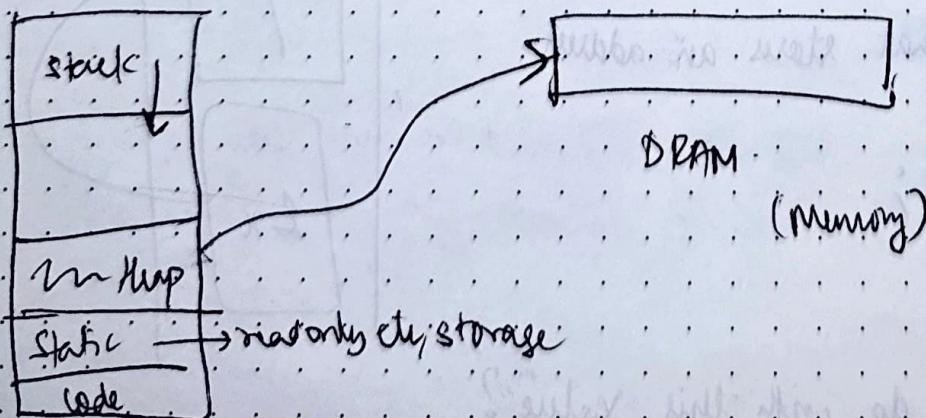
new / delete

## DYNAMIC MEMORY ALLOCATION

int x = 7 ;

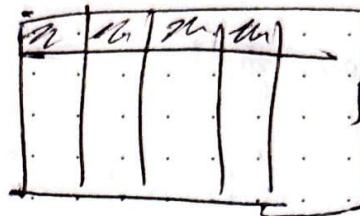
Block scope

If x is not available here  $\rightarrow$  as scope ends



If we want to allocate memory

int \*x = new int;  
↳ allocating memory  
↑  
\*mp



delete x;

int \*AllocateMemory() {

return new int;

}

int main() {

int \*x = AllocateMemory();

delete x;

return 0;

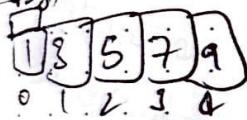
}

\* Allows us to allocate resources at runtime.

\* That resource is returned to us in a pointer.

delete[] student\_ids;

↳ deletes the array



int array[] = {1, 3, 5, 7, 9};

int array[5]

std::cout << array[0] << std::endl;

int \*px = array; // points to first index

cout << px;

gives memory of array.

discrepancy: \*px gives 1/ first element of the array)

px ++ → goes up by 4 bytes

\*px → gives 3

// Show the array offset

$\&(px+0)$      $\&(px+2)$      $\&(px+4)$  } → gives the full array  
 $\&(px+1)$      $\&(px+3)$

→ no need of a pointer for the above as

$\&(\text{array}+0)$      $\&(\text{array}+2)$      $\&(\text{array}+4)$  }

$\&(\text{array}+1)$      $\&(\text{array}+3)$

variable by default points to the beginning

array[0] { you are de-referencing the pointer

array[1] } It is the same as these

### Passing arrays into functions

```
void PrintArray(int arr[]){  
    for(int i=0; i <
```

It is a good practice to  
pass arrays with  
by reference  
to avoid  
copies.

}

delegating

= size(array)

size of (int)

int main(){

int array[] = {1, 3, 5, 7};

PrintArray(array);

return 0;

}

## Pitfalls of pointers

```
int main(){
```

int<sup>a</sup> Px = nullptr;

$$px = 42;$$

Zehm 07

segmentation fault, dependency

something which we do not have

also to

We forgot to de-allocate memory

white (brown)

int arry = new int[100];

return 0;

eventually there will be

memory test & we will

run out of memory.

## Dangling Pointers

That refers to a memory location that has already been freed or deleted

but the pointer still holds the address that location

do degeneration of that party can lead to undefined behavior.

```
5. int *ptr = new int(5); //Allocate memory
```

debt p<sup>t</sup>1

*x*      *ptr* = 20;

undifined

We are not restricted to NVII (or) multiplets

the point still holds the old address

int getPointer();

~~int localVar = 10;~~

return & location

Local rays in a fm. exist in

Stack memory is all destroyed

one—the first returns.

#### ④. double free

Deleting memory twice causes delete free.

Freeing the same pointer twice

```
int* ptr = int*(malloc(sizeof(int));  
free(ptr);  
free(ptr);  
new int;  
delete ptr;  
delete ptr;
```

#### → Function Pointers

```
int add(int x, int y) {  
    return x + y;  
}
```

```
int multiply(int x, int y) {  
    return x * y;  
}
```

```
int main() {  
    int (*op)(int, int);  
    // Name of fn pointer
```

op = add;

Cont < op(2, 2);

Button



Callback  
for

function pt.

#### → move in move semantics

Efficient transfer ~~form~~ of resources from 1 object to another object.

std::string s1 = "long string"; → value → no named storage

↳ Value

(or) location

↳ location where we are going to store

~~82 = \$1 ;~~

→ value → locator value.

↳ object that persists beyond a single expression. It has a memory address.

→ &value → temporary value that does not persist beyond the expression in which it is used. It is not associated with a memory address.

std::unique\_ptr

### SMART POINTERS

→ Raw pointers allows sharing

int x = 42;

int\* ptr = &x;

int\* ptr2 = ptr;

int\* x = new int[100];

for (int i=0; i<100; i++) {

\*x[i] = i;

}

apt = ~~new~~ X

→ later say we delete X

on or

& delete[] x;

apt = ~~new~~ X

→ gives 0

apt = ~~new~~ X

Smart pointers address these issues.

(RAII - scoped)

```
#include <memory> //unique_ptr
```

```
class VDT {
```

```
public:
```

```
    VDT() { std::cout << "Const created"; }
```

```
    ~VDT() { std::cout << "destruct" << std::endl; }
```

```
int main() {
```

```
    std::unique_ptr<VDT> mike = std::unique_ptr<VDT>(
```

```
        new VDT());
```

We don't have to think about deleting the memory created.

//create an "array", that is pointed to by one unique pointer

```
std::unique_ptr<VDT[]> mike_array = std::unique_ptr<VDT[]>(
```

(new VDT[10]);

↑

The equivalent to the line above is

```
std::unique_ptr<VDT[]> mike_array = std::make_unique(VDT[])(10);
```

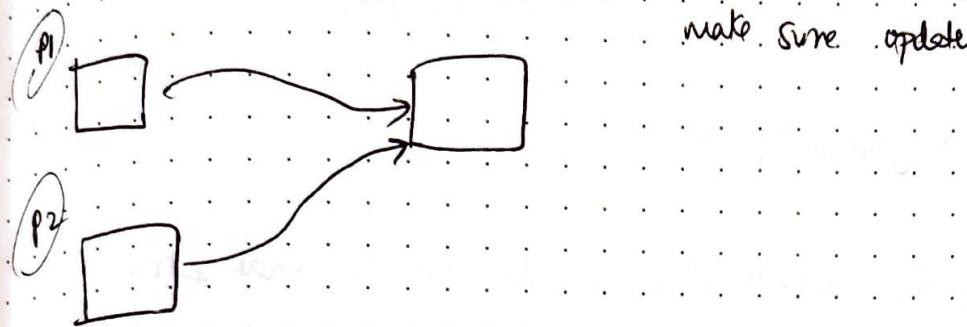
Factory Pattern

→ 2. Unique pointers cannot point to the same box of memory

## Unique pointer

- (1) Cannot Share
- (2) No copies
- (3) No worries about deleting  
(Prevents memory leak)
- (4) Move is allowed
- (5) Simplifies memory management.

→ shared\_ptr



make sure updates

→ unique\_ptr was a scoped pointer

→ shared\_ptr is a smart pointer that retains shared ownership of an object through pointers

The object will be destroyed only when the last shared\_ptr that points-to it is destroyed. (or) rest

make\_shared !      use\_count()

shared\_ptr<int> ptr1 =

std::make\_shared<int>(42);

To count no. of pointers are

pointing



std::shared\_ptr<int> ptr2 = ptr1;

std::cout << \*ptr2 // output 42

ptr1.use\_count() → 2

as block scope ended

g

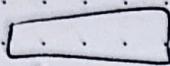
ptr1.use\_count(1 → gives 1)

## Weak\_ptr

+ smart pointer (a non-counting)

- Q) ~~forgetting~~ acting smart pointer but holds a non-counting pointer, does not increase the reference count.

ptr 1



ptr 2



free ptr 1

ptr2[0] → gives undefined

(dangling ptr)

So, definitely if this is still allowed when we use weak\_ptr -

~~std::weak\_ptr<VDT> ptr2 = ptr1;~~

std::shared\_ptr<VDT> ptr1 = std::make\_shared<VDT>();  
{

std::::weak\_ptr<VDT> ptr2 = ptr1

ptr2.use\_count() → still gives 1

y

ptr1.use\_count() → gives 1

Interface (.hpp or .h) and Implementation (.cpp, .cxx, .c)

#include "mikemath.hpp" → our local path

#include <iostream> → present in some system file path

class & struct

C